

Linux Kernel Introduction

...

WHAT TO EXPECT

HISTORY

LINUX KERNEL vs LINUX OS

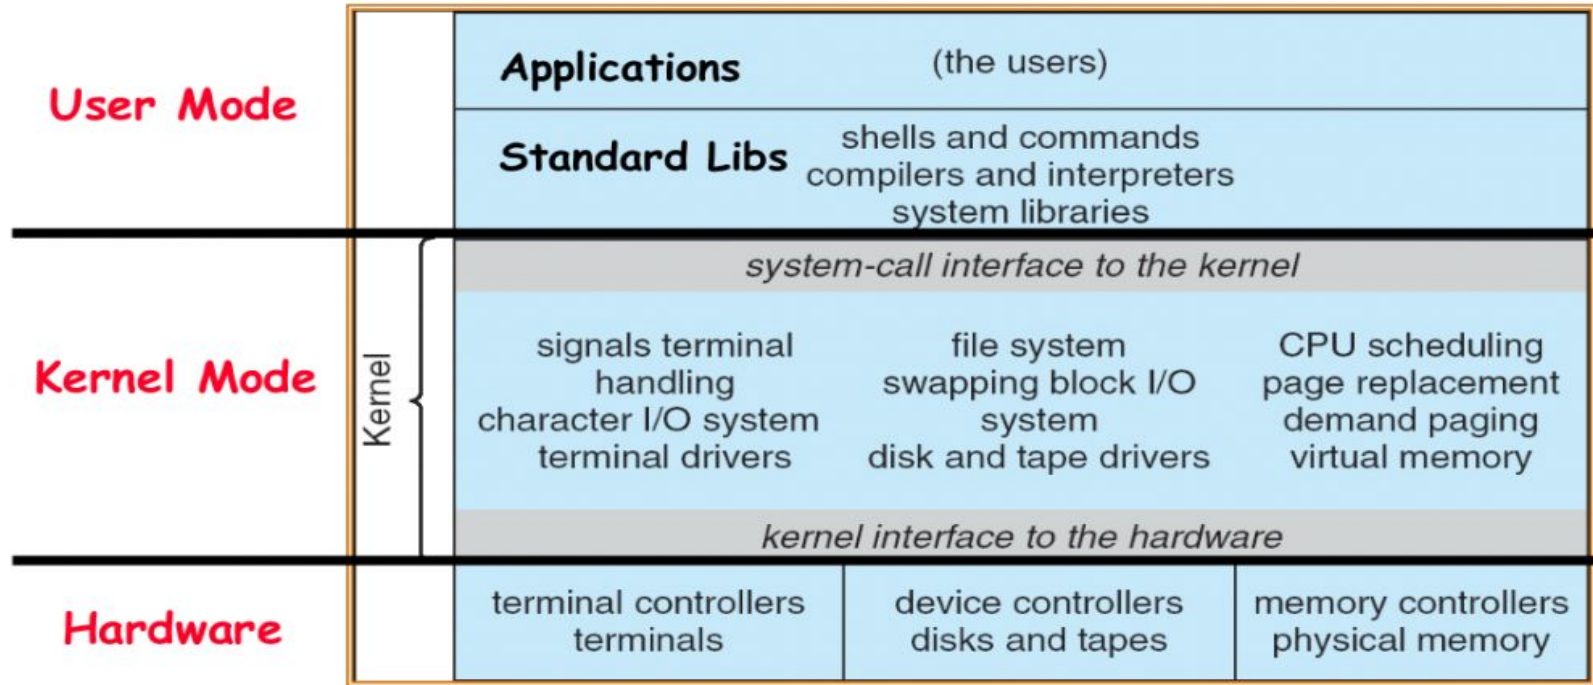


Image taken from: <https://www.linuxfordevices.com/tutorials/linux/linux-kernel>

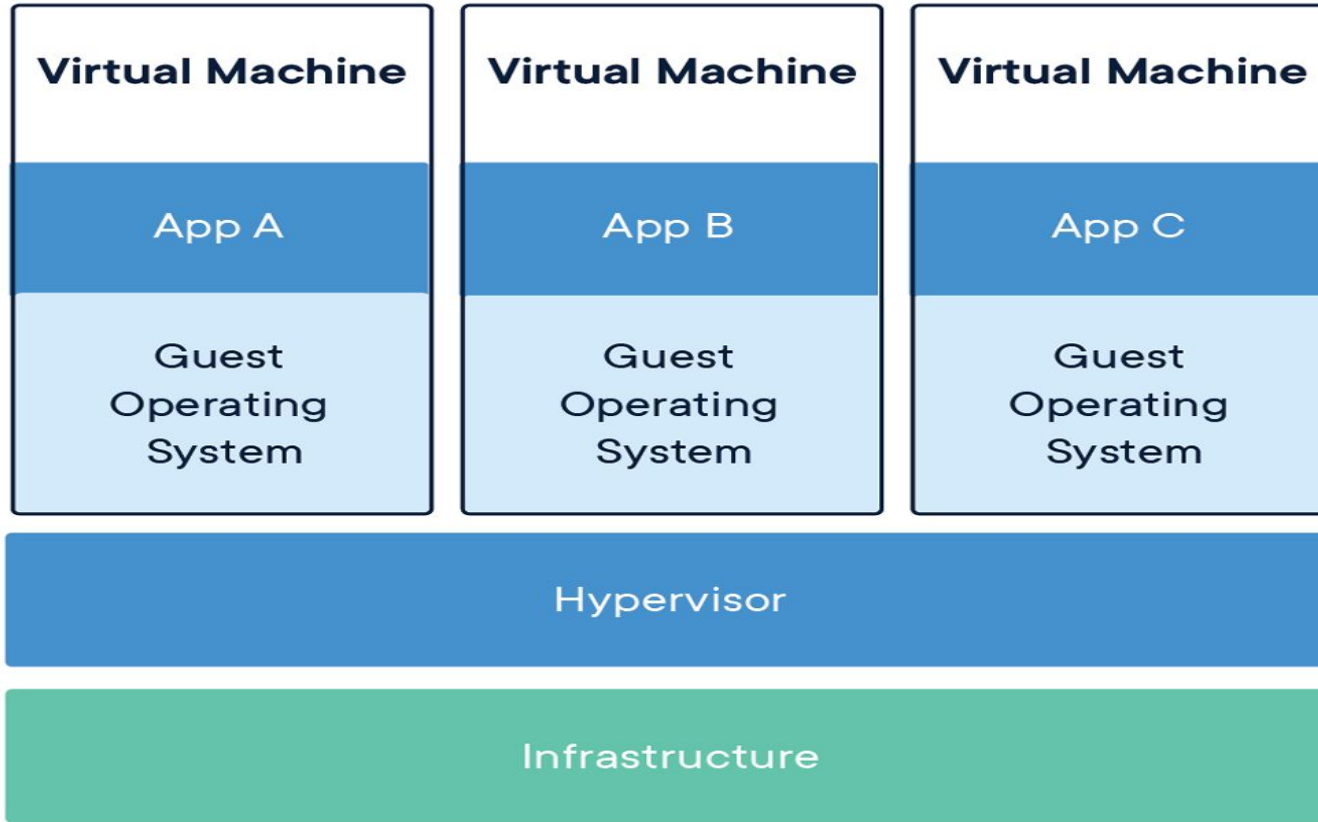


Image taken from: https://commons.wikimedia.org/wiki/File:Container-vm-whatcontainer_2.png

LINUX KERNEL IS

THE open source project

Acknowledgements

While the Linux project has been closely associated with me personally, partly due to the name, I would like to make it very clear that the Linux operating system is a huge project done co-operatively by lots of people all over the world.

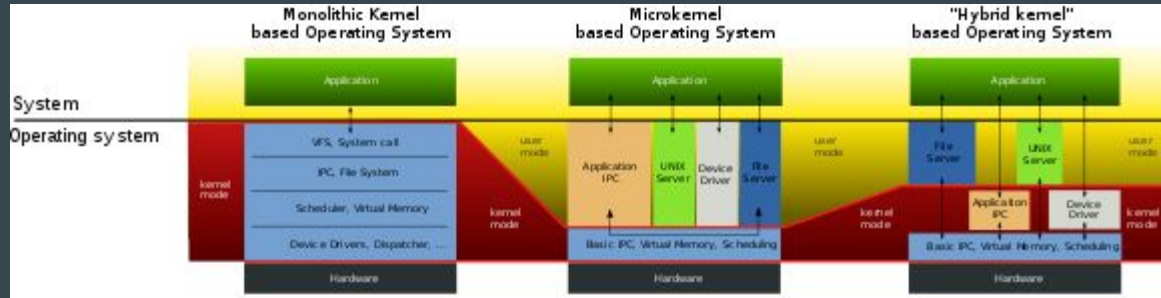
Even if you discount all the user-level programs that are an integral part of any running Linux system, just the kernel contains code from hundreds of people from all around the world.

Thanks to all of you.

https://www.cs.helsinki.fi/u/kutvonen/index_files/linus.pdf

UNIX (kinda)

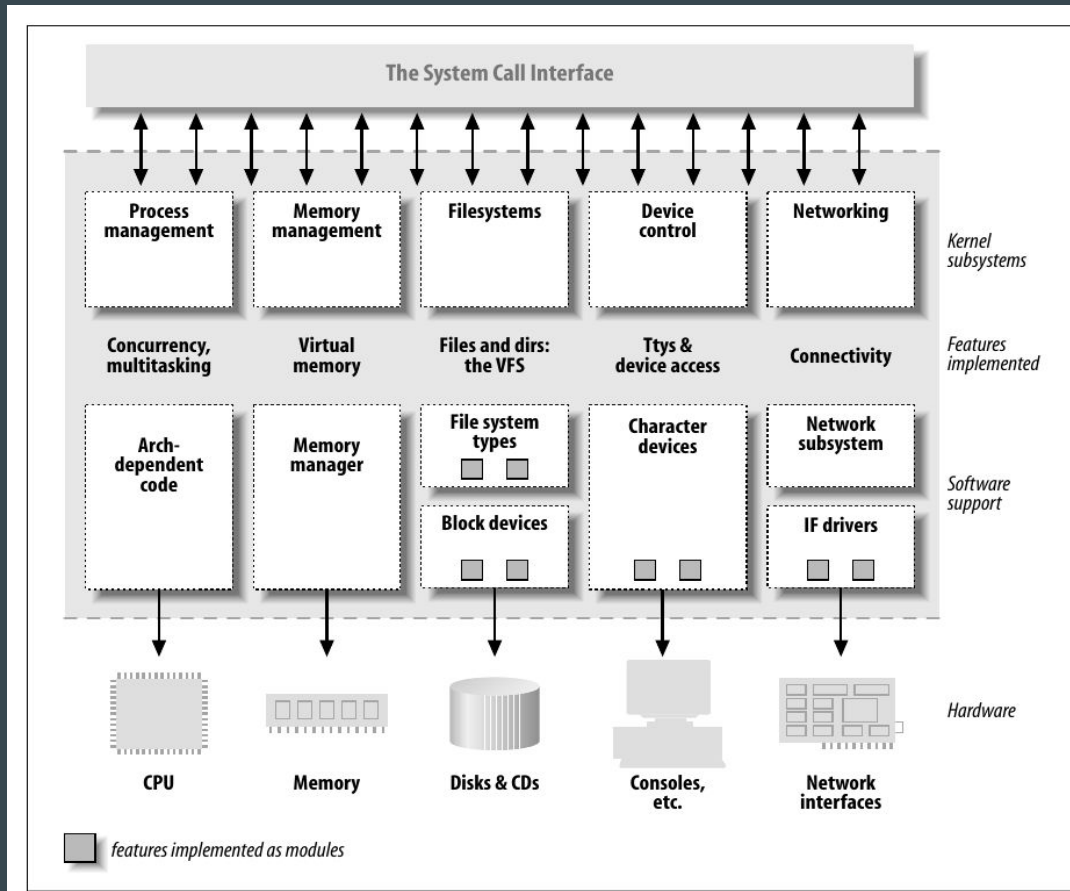
Monolithic...



Taken from:

<https://upload.wikimedia.org/wikipedia/commons/thumb/d/d0/OS-structure2.svg/580px-OS-structure2.svg.png>

... but with Modules



Taken from: <https://lwn.net/Kernel/LDD3/>

Symmetric Multiprocessors

Preemptive


Reentrant

HOW LINUX WORKS


Process vs Thread

~~Process vs Thread~~ Task

(Lord of the) Rings



Least privileged



Most privileged

Traps

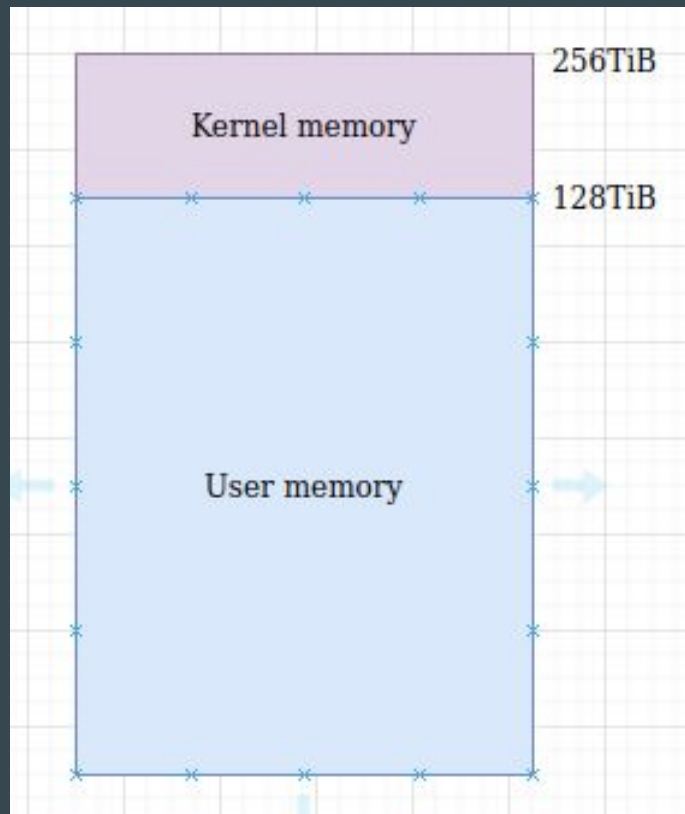
Cutest traps in anime

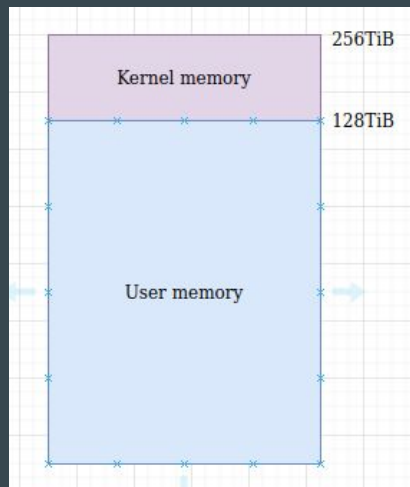
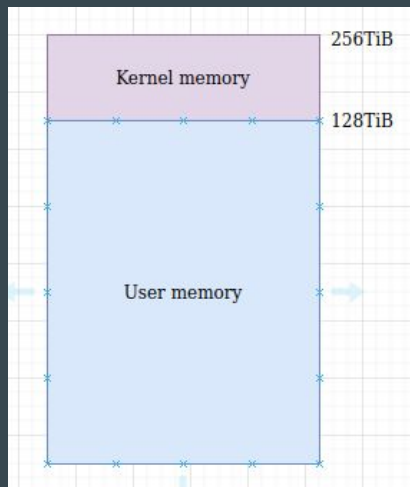
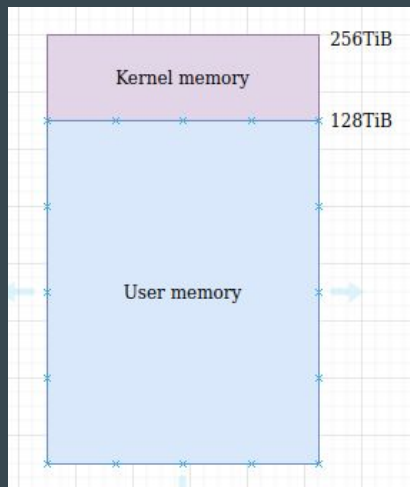
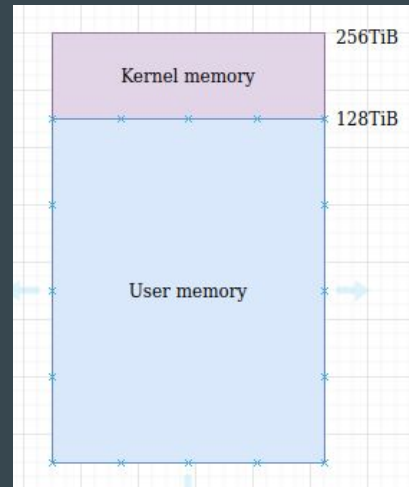
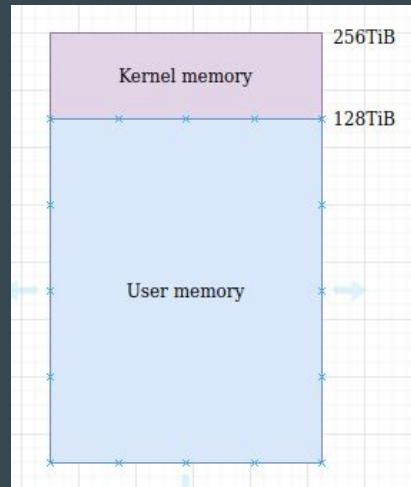
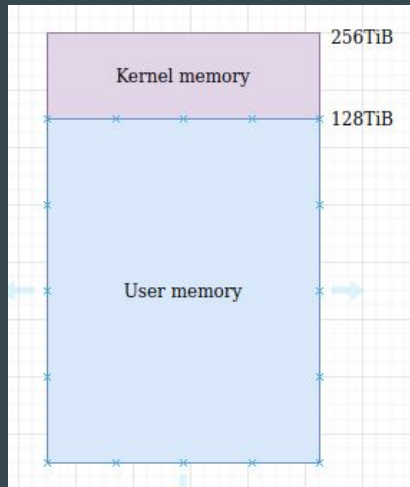
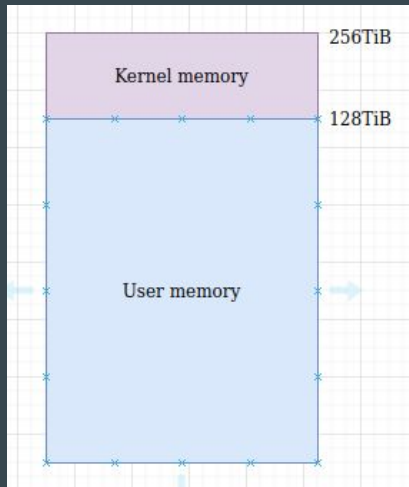


Scheduling...

**... Or how you can write on the terminal without
freezing your entire machine.**

Userland vs Kernelland

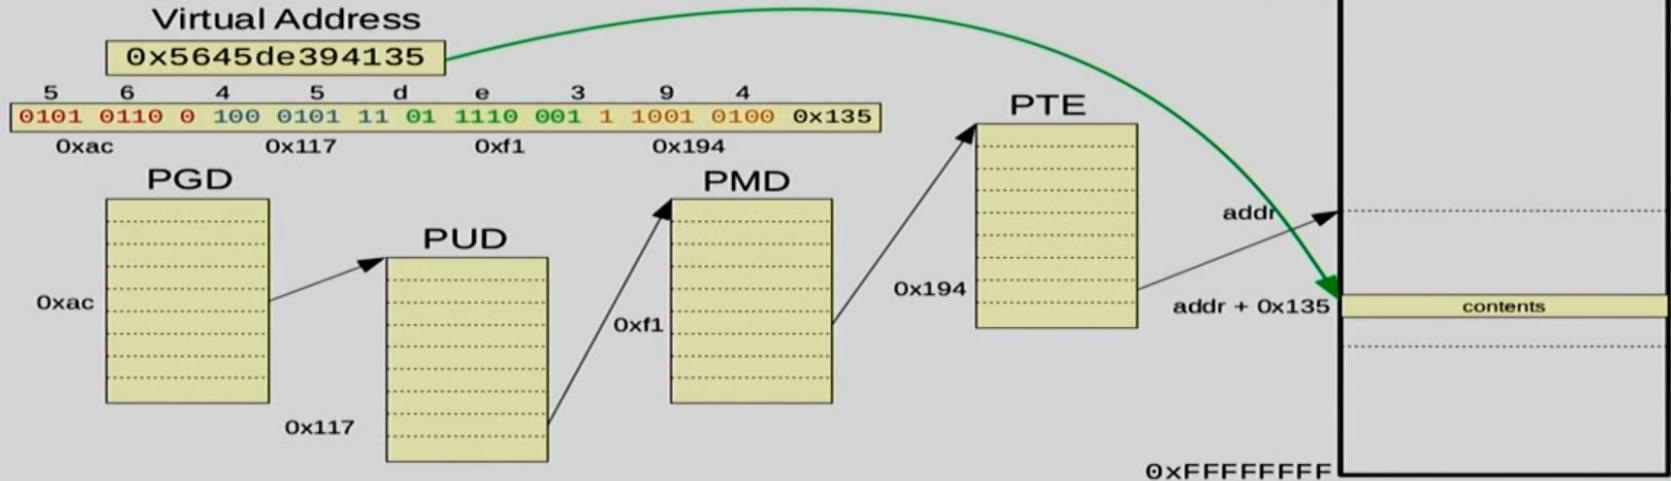




?????

Everything in OS can be solved with one more layer
of indirection

Page Tables



vmware

Steven Rostedt - Learning the Linux Kernel with tracing

<https://www.youtube.com/watch?v=JRyrhsx-L5Y&t=3128s>

HOW TO TALK TO THE KERNEL

Syscalls

HACKING THE KERNEL

No libc

No floats

(almost) No stack

No swapping

No userland memory

Lots of spinlocks

Oops & Panic

```
printk(KERN_ALERT "Hello, world\n");
```

kalloc()

`copy_from_user()`
`copy_to_user()`

**Clone Torvalds tree
or
Download the sources from kernel.org**

elixir.bootlin.com

COMPILING

ANATOMY OF A BUG

Capabilities

```
(base) → linux git:(main) lsb_release -a
No LSB modules are available.
Distributor ID: LinuxMint
Description:    Linux Mint 19.3 Tricia
Release:        19.3
Codename:       tricia
(base) → linux git:(main) ls -l /bin/ping
-rwsr-xr-x 1 root root 64424 Jun 28  2019 /bin/ping
(base) → linux git:(main) █
```

VS

```
(base) vccolombo@laptop:~$ lsb_release -a
No LSB modules are available.
Distributor ID: Neon
Description:    KDE neon User Edition 5.22
Release:        20.04
Codename:       focal
(base) vccolombo@laptop:~$ ls -l /bin/ping
-rwxr-xr-x 1 root root 72776 jan 30  2020 /bin/ping
(base) vccolombo@laptop:~$ █
```

CVE-2016-9793

CVE-2016-9793

The `sock_setsockopt` function in `net/core/sock.c` in the Linux kernel before 4.8.14 mishandles negative values of `sk_sndbuf` and `sk_rcvbuf`, which allows local users to cause a denial of service (memory corruption and system crash) or possibly have unspecified other impact by leveraging the `CAP_NET_ADMIN` capability for a crafted `setsockopt` system call with the (1) `SO_SNDBUFFORCE` or (2) `SO_RCVBUFFORCE` option.

WHERE TO GO FROM HERE

Working with the Kernel

Contributing to the Kernel

Exploiting the Kernel

```
printk(KERN_NOTICE "Thanks for watching");
```